

# Neural Processes

Alexander Plakhin, Nikolai Averianov

HSE FES Probability Theory Club

November 11, 2023

- Gaussian Processes Regression provides a probabilistic framework to impose prior distribution over a set of functions and update this distribution after observing the data
- Obvious drawbacks are computational complexity ( $\mathcal{O}(N^3)$ ) and limited expressiveness
- Neural Processes is a similar framework which addresses the issues above

# Gaussian Processes Regression (Reminder)

- We assume that you know what is a Gaussian Processes Regression

- Let's consider the process  $F : \mathcal{X} \rightarrow \mathcal{Y}$  as a random function. We define distribution over  $Y_{1:n} := (F(x_1), \dots, F(x_n))$
- Imagine we have a collection of distributions  $\rho_{x_{1:n}}$ . What are the conditions for this collection to define a stochastic process  $F$ ?
- Exchangeability (for any permutation):

$$\rho_{x_1, \dots, x_n}(y_1, \dots, y_n) = \rho_{x_{\pi(1)}, \dots, x_{\pi(n)}}(y_{\pi(1)}, \dots, y_{\pi(n)})$$

- Consistency:

$$\rho_{x_{1:m}}(y_{1:m}) = \int \rho_{x_{1:n}}(y_{1:n}) dy_{m+1:n}$$

- These conditions are sufficient by the Kolmogorov Extension Theorem
- We need to build a Neural Network that satisfies the conditions above

# Model Setup

- Given a particular instantiation of the stochastic process  $f$ :

$$\rho_{x_{1:n}}(y_{1:n}) = \int p(y_{1:n}|f, x_{1:n})p(f)df$$

- if we add noise:  $Y_i \sim \mathcal{N}(F(x_i), \sigma^2)$ :

$$p(y_{1:n}|x_{1:n}) = \int p(f) \prod_{i=1}^n \mathcal{N}(y_i|f(x_i), \sigma^2)df$$

- Now we can do the following: let's assume that realization of the process depends on a global latent variable  $z$  and parameterize  $F(x)$  as a Neural Network  $g_\theta(x, z)$ :

$$p(z, y_{1:n}|x_{1:n}) = p(z) \prod_{i=1}^n \mathcal{N}(y_i|g_\theta(x_i, z), \sigma^2)$$

- Let  $q$  be variational posterior parametrized by another NN. ELBO (expectations are w.r.t  $q(z|x_{1:n}, y_{1:n})$ ):

$$\log p(y_{1:n}|x_{1:n}) \geq \mathbb{E} \left[ \sum_{i=1}^n \log p(y_i|z, x_i) + \log \frac{p(z)}{q(z|x_{1:n}, y_{1:n})} \right]$$

- But for this problem we really care more about conditional sampling. So we have:

$$\log p(y_{m+1:n}|x_{1:n}, y_{1:m}) \geq \mathbb{E} \left[ \sum_{i=m+1}^n \log p(y_i|z, x_i) + \log \frac{p(z|x_{1:m}, y_{1:m})}{q(z|x_{1:n}, y_{1:n})} \right]$$


- Of course we have no clue about  $p(z|x_{1:n}, y_{1:m}) \Rightarrow$  we end up maximizing:

$$\log p(y_{m+1:n}|x_{1:n}, y_{1:m}) \geq \mathbb{E} \left[ \sum_{i=m+1}^n \log p(y_i|z, x_i) + \log \frac{q(z|x_{1:m}, y_{1:m})}{q(z|x_{1:n}, y_{1:n})} \right]$$


# Architecture

The model consist of several parts:

- Encoder  $h$  that takes pairs  $(x, y)_i$  and maps them to representation  $r_i$
- Aggregator  $a$  that summarizes inputs and is order-invariant. We need a single global representation  $r = a(r_1, \dots, r_n)$  to parameterize  $z \sim \mathcal{N}(\mu(r), \sigma^2(r)I)$ . Note linear complexity of aggregator in case  $a(r_1, \dots, r_n) = \sum_{i=1}^n r_i$
- Conditional decoder  $g$  that takes new points  $x$  and the sampled latent variable  $z$




(a) Graphical model




(b) Computational diagram


# Similar models




(a) Conditional VAE



(b) Neural statistician







(c) Conditional neural process



(d) Neural process



# Results



- 1 Neural Processes – <https://arxiv.org/pdf/1807.01622.pdf>
- 2 <https://yanndubs.github.io/Neural-Process-Family/text/Intro.html>
- 3 <https://github.com/EmilienDupont/neural-processes>